

A Transition-Based Directed Acyclic Graph Parser for Universal Conceptual Cognitive Annotation

Daniel Hershcovich, Omri Abend and Ari Rappoport




ACL 2017

TUPA — Transition-based UCCA Parser

The first parser to support the combination of

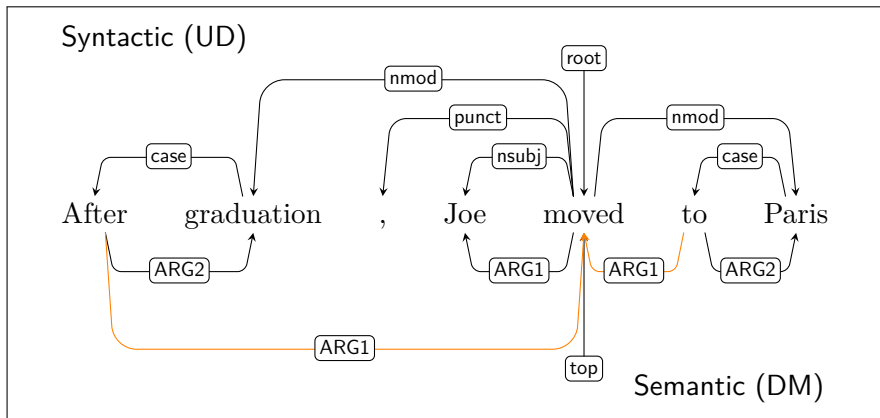
1. Non-terminal nodes
2. Reentrancy
3. Discontinuity

— needed for many semantic schemes (e.g. AMR, UCCA).



Linguistic Structure Annotation Schemes

- Syntactic dependencies (Nivre, 2005)
- Semantic dependencies (Oepen et al., 2016)



Bilexical dependencies.

Linguistic Structure Annotation Schemes


- Syntactic dependencies (Nivre, 2005)
- Semantic dependencies (Oepen et al., 2016)
- AMR (Banarescu et al., 2013)
- UCCA (Abend and Rappoport, 2013)
- Other semantic representation schemes¹

...showered = ...took a shower

...’s war against crime = ...fights crime

¹See recent survey (Abend and Rappoport, 2017)

Universal Conceptual Cognitive Annotation (UCCA)



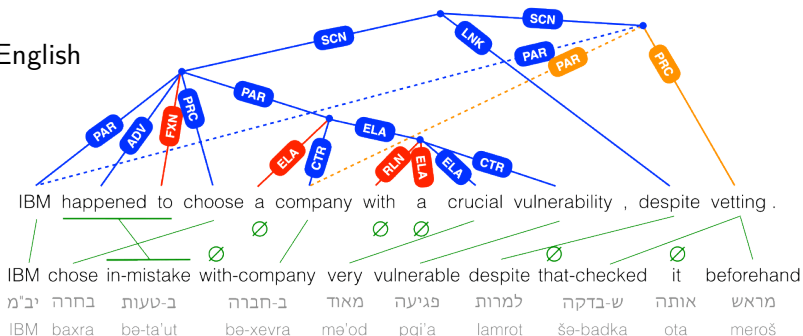
After graduation, Joe moved to Paris

<i>P</i> process	<i>S</i> state	<i>A</i> participant
<i>L</i> linker	<i>H</i> linked scene	<i>C</i> center
<i>E</i> elaborator	<i>D</i> adverbial	<i>R</i> relator
<i>N</i> connector	<i>U</i> punctuation	<i>F</i> function
<i>G</i> ground		


The UCCA Semantic Representation Scheme

Cross-linguistically applicable (Abend and Rappoport, 2013).
Stable in translation (Sulem et al., 2015).

English



Hebrew



UCCAApp

Rapid and intuitive annotation interface (Abend et al., 2017).
Usable by non-experts. <http://ucca-demo.cs.huji.ac.il>

Linker (L)	i
Ground (G)	i
Participant (A)	i
State (S)	i
Process (P)	i
Adverbial (D)	i
Time (T)	i
Center (C)	i
Elaborator (E)	i
Connector (N)	i
Relator (R)	i
Uncertain (UNC)	i
Unanalyzable (UI)	i
Function (F)	i

William Bradley Pitt was born in Shawnee, Oklahoma , to William Alvin Pitt , who ran a trucking company , and Jane Etta (née Hillhouse) , a school counsellor . The family soon moved to Springfield , Missouri , where he lived together with his younger siblings , Douglas (born 1966) and Julie Neal (born 1969). Born into a conservative household , he was raised as Southern Baptist , but has since stated that he does not " have a great relationship with religion " and that he " oscillates between agnosticism and atheism . " Pitt has described Springfield as " Mark Twain country " , Jesse James country " , having grown up with " a lot of hills , a lot of lakes " .

1 H William Bradley Pitt was born in Shawnee , Oklahoma + F x

1-1 A William Bradley Pitt + F x

1-2 F was + F x

1-3 P born + F x

1-4 A in Shawnee , Oklahoma + F x

1-4-1 R in + F x

1-4-2 C UNA Shawnee , Oklahoma + F x

HUME

UCCA facilitates semantics-based human evaluation of machine translation (Birch et al., 2016). <http://ucca.cs.huji.ac.il/mteval>

Für leicht fettleibige Diabetiker kann die Gewichtsschaden - OP hilfreich sein

A B ● ● ● ● ×

For mildly obese diabetics, weight loss surgery may be helpful


■ For mildly obese diabetics, weight loss surgery may be helpful Für leicht fettleibige Diabetiker kann die Gewichtsschaden - OP hilfreich sein A B ● ● ● ● ×

- For mildly obese diabetics Für leicht fettleibige Diabetiker kann A B ● ● ● ● ×
 - For Für A B ● ● ● ● ×
 - mildly leicht A B ● ● ● ● ×
 - obese fettleibige A B ● ● ● ● ×
 - diabetics Diabetiker kann A B ● ● ● ● ×
- weight loss surgery Gewichtsschaden - A B ● ● ● ● ×
 - weight loss Gewichtsschaden A B ● ● ● ● ×
 - weight Gewichtsschaden A B ● ● ● ● ×
 - loss A B ● ● ● ● ×
 - surgery - A B ● ● ● ● ×
- may A B ● ● ● ● ×
- be helpful hilfreich sein A B ● ● ● ● ×
 - be sein A B ● ● ● ● ×
 - helpful hilfreich sein A B ● ● ● ● ×

Graph Structure

UCCA generates a directed acyclic graph (DAG), where the text tokens are terminals. Structural properties:

1. **Non-terminal nodes** — hierarchy of entities and events




You want to take a long bath

Graph Structure

UCCA generates a directed acyclic graph (DAG), where the text tokens are terminals. Structural properties:

1. **Non-terminal nodes** — hierarchy of entities and events
2. **Reentrancy** — allow argument sharing




You want to take a long bath

Graph Structure

UCCA generates a directed acyclic graph (DAG), where the text tokens are terminals. Structural properties:

1. **Non-terminal nodes** — hierarchy of entities and events
2. **Reentrancy** — allow argument sharing
3. **Discontinuity** – conceptual units are split



You want to take a long bath

Transition-Based Parsing

First used for dependency parsing (Nivre, 2004).

Parse text $w_1 \dots w_n$ to graph $G = (V, E, \ell)$ incrementally.

Transition-Based Parsing

First used for dependency parsing (Nivre, 2004).

Parse text $w_1 \dots w_n$ to graph $G = (V, E, \ell)$ incrementally.

Initial state:

stack

buffer



You	want	to	take	a	long	bath
-----	------	----	------	---	------	------

Transition-Based Parsing

First used for dependency parsing (Nivre, 2004).

Parse text $w_1 \dots w_n$ to graph $G = (V, E, \ell)$ incrementally.

Initial state:

stack

buffer



You	want	to	take	a	long	bath
-----	------	----	------	---	------	------

TUPA transitions:

$\{\text{SHIFT}, \text{REDUCE}, \text{NODE}_X, \text{LEFT-EDGE}_X, \text{RIGHT-EDGE}_X, \\ \text{LEFT-REMOTE}_X, \text{RIGHT-REMOTE}_X, \text{SWAP}, \text{FINISH}\}$

Support **non-terminal nodes**, **reentrancy** and **discontinuity**.

Transition-Based UCCA Parsing

⇒ SHIFT

stack

buffer

●	You
---	-----

want	to	take	a	long	bath
------	----	------	---	------	------

graph



Transition-Based UCCA Parsing


\Rightarrow RIGHT-EDGE_A

stack

buffer

●	You
---	-----

want	to	take	a	long	bath
------	----	------	---	------	------



Transition-Based UCCA Parsing

⇒ SHIFT


stack

buffer

●	You	want
---	-----	------

to	take	a	long	bath
----	------	---	------	------

graph



Transition-Based UCCA Parsing

⇒ SWAP


stack

buffer

●	want
---	------

You	to	take	a	long	bath
-----	----	------	---	------	------

graph






Transition-Based UCCA Parsing

\Rightarrow RIGHT-EDGE_P

stack

buffer




Transition-Based UCCA Parsing

⇒ REDUCE


stack



buffer



graph





Transition-Based UCCA Parsing


\Rightarrow SHIFT

stack

buffer



graph



Transition-Based UCCA Parsing

⇒ SHIFT


stack

buffer

●	You	to
---	-----	----

take	a	long	bath
------	---	------	------

graph



Transition-Based UCCA Parsing

$\Rightarrow \text{NODE}_F$


stack

buffer

●	You	to
---	-----	----

●	take	a	long	bath
---	------	---	------	------

graph



Transition-Based UCCA Parsing


⇒ REDUCE

stack

buffer

●	You
---	-----

●	take	a	long	bath
---	------	---	------	------





Transition-Based UCCA Parsing


⇒ SHIFT

stack

buffer



graph



Transition-Based UCCA Parsing


⇒ SHIFT

stack

buffer

●	You	●	take
---	-----	---	------

a	long	bath
---	------	------



Transition-Based UCCA Parsing


$\Rightarrow \text{NODE}_C$

stack

buffer

●	You	●	take
---	-----	---	------

●	a	long	bath
---	---	------	------






Transition-Based UCCA Parsing

⇒ REDUCE

stack

buffer





Transition-Based UCCA Parsing

⇒ SHIFT

stack

buffer




Transition-Based UCCA Parsing

\Rightarrow RIGHT-EDGE_P

stack

buffer




Transition-Based UCCA Parsing

⇒ SHIFT


stack



buffer



graph




Transition-Based UCCA Parsing

\Rightarrow RIGHT-EDGE_F

stack


buffer




Transition-Based UCCA Parsing

⇒ REDUCE


stack



buffer



graph



Transition-Based UCCA Parsing

⇒ SHIFT


stack



buffer



graph





Transition-Based UCCA Parsing

⇒ SWAP

stack

buffer




Transition-Based UCCA Parsing

\Rightarrow RIGHT-EDGE_D

stack


buffer




Transition-Based UCCA Parsing

⇒ REDUCE

stack



buffer




Transition-Based UCCA Parsing

⇒ SWAP

stack



buffer




Transition-Based UCCA Parsing

\Rightarrow RIGHT-EDGE_A

stack



buffer




Transition-Based UCCA Parsing

⇒ REDUCE

stack



buffer




Transition-Based UCCA Parsing

⇒ REDUCE

stack



buffer



Transition-Based UCCA Parsing

⇒ SHIFT

stack


You

buffer



bath

graph



Transition-Based UCCA Parsing


⇒ SHIFT

stack

You ●

buffer

bath




Transition-Based UCCA Parsing

\Rightarrow LEFT-REMOTE_A

stack



buffer




Transition-Based UCCA Parsing

⇒ SHIFT

stack

buffer




Transition-Based UCCA Parsing

\Rightarrow RIGHT-EDGE_C

stack

buffer




Transition-Based UCCA Parsing

⇒ FINISH

stack

buffer

You	●	bath
-----	---	------



TUPA Model

Greedy parsing, experimenting with three classifiers:

- Sparse** Perceptron with sparse features.
- MLP** Embeddings + feedforward NN classifier.
- BiLSTM** Embeddings + deep bidirectional LSTM + MLP
(Kiperwasser and Goldberg, 2016).

Features: words, POS, syntactic dependencies, existing edge labels from the stack and buffer + parents, children, grandchildren; ordinal features (height, number of parents and children)


stack




buffer




BiLSTM




BiLSTM




BiLSTM




BiLSTM





stack



buffer




graph



Training


An *oracle* provides the transition sequence given the correct graph:



SHIFT, RIGHT-EDGE_A, SHIFT, SWAP, RIGHT-EDGE_P, REDUCE, SHIFT,
 SHIFT, NODE_F, REDUCE, SHIFT, SHIFT, NODE_C, REDUCE, SHIFT,
 RIGHT-EDGE_P, SHIFT, RIGHT-EDGE_F, REDUCE, SHIFT, SWAP,
 RIGHT-EDGE_D, REDUCE, SWAP, RIGHT-EDGE_A, REDUCE, REDUCE, SHIFT,
 SHIFT, LEFT-REMOTE_A, SHIFT, RIGHT-EDGE_C, FINISH

Experimental Setup

- UCCA Wikipedia corpus ($4268 + 454 + 503$ sentences).
- Out-of-domain: English part of English-French parallel corpus, *Twenty Thousand Leagues Under the Sea* (506 sentences).




Baselines

Bilexical DAG parsers (allow **reentrancy**):

- DAGParser (Ribeyre et al., 2014): transition-based.
- TurboParser (Almeida and Martins, 2015): graph-based.

Tree parsers (all transition-based):

- MaltParser (Nivre et al., 2007): bilexical tree parser.
- LSTM Parser (Dyer et al., 2015): bilexical tree parser, NN.
- UPARSE (Maier, 2015): allows **non-terminals**, **discontinuity**.



Bilexical DAG approximation (for tree, delete remote edges).

Results

TUPA_{BiLSTM} obtains the highest F-scores in all metrics:

	Primary			Remote		
	LP	LR	LF	LP	LR	LF
TUPA _{Sparse}	64.5	63.7	64.1	19.8	13.4	16
TUPA _{MLP}	65.2	64.6	64.9	23.7	13.2	16.9
TUPA _{BiLSTM}	74.4	72.7	73.5	47.4	51.6	49.4
Bilexical DAG	(91)			(58.3)		
DAGParser	61.8	55.8	58.6	9.5	0.5	1
TurboParser	57.7	46	51.2	77.8	1.8	3.7
Bilexical tree	(91)			—		
MaltParser	62.8	57.7	60.2	—	—	—
LSTM Parser	73.2	66.9	69.9	—	—	—
Tree	(100)			—		
UPARSE	60.9	61.2	61.1	—	—	—

Results on the Wiki test set.

Results

Comparable out-of-domain test set:

	Primary			Remote		
	LP	LR	LF	LP	LR	LF
TUPA _{Sparse}	59.6	59.9	59.8	22.2	7.7	11.5
TUPA _{MLP}	62.3	62.6	62.5	20.9	6.3	9.7
TUPA _{BiLSTM}	68.7	68.5	68.6	38.6	18.8	25.3
Bilexical DAG			(91.3)			(43.4)
DAGParser	56.4	50.6	53.4	—	0	0
TurboParser	50.3	37.7	43.1	100	0.4	0.8
Bilexical tree			(91.3)			—
MaltParser	57.8	53	55.3	—	—	—
LSTM Parser	66.1	61.1	63.5	—	—	—
Tree			(100)			—
UPARSE	52.7	52.8	52.8	—	—	—

Results on the 20K Leagues out-of-domain set.

Conclusion

- UCCA's semantic distinctions require a graph structure including **non-terminals**, **reentrancy** and **discontinuity**.
- TUPA is an accurate transition-based UCCA parser, and the first to support any DAG over the text's tokens.
- Outperforms strong conversion-based baselines.

Code: <https://github.com/danielhers/tupa>

Demo: <http://bit.ly/tupademo>

Corpora: <http://www.cs.huji.ac.il/~oabend/ucca.html>

Conclusion

- UCCA's semantic distinctions require a graph structure including **non-terminals**, **reentrancy** and **discontinuity**.
- TUPA is an accurate transition-based UCCA parser, and the first to support any DAG over the text's tokens.
- Outperforms strong conversion-based baselines.

Future Work:

- More languages (German corpus construction is underway).
- Parsing other schemes, such as AMR.
- Text simplification, MT evaluation and other applications.

Code: <https://github.com/danielhers/tupa>

Demo: <http://bit.ly/tupademo>

Corpora: <http://www.cs.huji.ac.il/~oabend/ucca.html>

Thank you

References I

- Abend, O. and Rappoport, A. (2013).
Universal Conceptual Cognitive Annotation (UCCA).
In *Proc. of ACL*, pages 228–238.
- Abend, O. and Rappoport, A. (2017).
The state of the art in semantic representation.
In *Proc. of ACL*.
to appear.
- Abend, O., Yerushalmi, S., and Rappoport, A. (2017).
UCCAApp: Web-application for syntactic and semantic phrase-based annotation.
In *Proc. of ACL: System Demonstration Papers*.
to appear.
- Almeida, M. S. C. and Martins, A. F. T. (2015).
Lisbon: Evaluating TurboSemanticParser on multiple languages and out-of-domain data.
In *Proc. of SemEval*, pages 970–973.
- Banarescu, L., Bonial, C., Cai, S., Georgescu, M., Griffitt, K., Hermjakob, U., Knight, K., Palmer, M., and Schneider, N. (2013).
Abstract Meaning Representation for sembanking.
In *Proc. of the Linguistic Annotation Workshop*.
- Birch, A., Abend, O., Bojar, O., and Haddow, B. (2016).
HUME: Human UCCA-based evaluation of machine translation.
In *Proc. of EMNLP*, pages 1264–1274.
- Dyer, C., Ballesteros, M., Ling, W., Matthews, A., and Smith, N. A. (2015).
Transition-based dependency parsing with stack long short-term memory.
In *Proc. of ACL*, pages 334–343.

References II

- Kiperwasser, E. and Goldberg, Y. (2016).
Simple and accurate dependency parsing using bidirectional LSTM feature representations.
TACL, 4:313–327.
- Maier, W. (2015).
Discontinuous incremental shift-reduce parsing.
In *Proc. of ACL*, pages 1202–1212.
- Nivre, J. (2004).
Incrementality in deterministic dependency parsing.
In Keller, F., Clark, S., Crocker, M., and Steedman, M., editors, *Proceedings of the ACL Workshop Incremental Parsing: Bringing Engineering and Cognition Together*, pages 50–57, Barcelona, Spain. Association for Computational Linguistics.
- Nivre, J. (2005).
Dependency grammar and dependency parsing.
Technical Report MSI 05133, Växjö University, School of Mathematics and Systems Engineering.
- Nivre, J., Hall, J., Nilsson, J., Chanev, A., Eryigit, G., Kübler, S., Marinov, S., and Marsi, E. (2007).
MaltParser: A language-independent system for data-driven dependency parsing.
Natural Language Engineering, 13(02):95–135.
- Oepen, S., Kuhlmann, M., Miyao, Y., Zeman, D., Cinková, S., Flickinger, D., Hajic, J., Ivanova, A., and Uresová, Z. (2016).
Towards comparability of linguistic graph banks for semantic parsing.
In *LREC*.
- Ribeyre, C., Villemonte de la Clergerie, E., and Seddah, D. (2014).
Alpage: Transition-based semantic graph parsing with syntactic features.
In *Proc. of SemEval*, pages 97–103.

References III

Sulem, E., Abend, O., and Rappoport, A. (2015).

Conceptual annotations preserve structure across translations: A French-English case study.

In *Proc. of S2MT*, pages 11–22.

UCCA Corpora

	Wiki			20K
	Train	Dev	Test	Leagues
# passages	300	34	33	154
# sentences	4268	454	503	506
# nodes	298,993	33,704	35,718	29,315
% terminal	42.96	43.54	42.87	42.09
% non-term.	58.33	57.60	58.35	60.01
% discont.	0.54	0.53	0.44	0.81
% reentrant	2.38	1.88	2.15	2.03
# edges	287,914	32,460	34,336	27,749
% primary	98.25	98.75	98.74	97.73
% remote	1.75	1.25	1.26	2.27
Average per non-terminal node				
# children	1.67	1.68	1.66	1.61

Corpus statistics.

Evaluation

Mutual edges between predicted graph $G_p = (V_p, E_p, \ell_p)$ and gold graph $G_g = (V_g, E_g, \ell_g)$, both over terminals $W = \{w_1, \dots, w_n\}$:

$$M(G_p, G_g) = \left\{ (e_1, e_2) \in E_p \times E_g \mid y(e_1) = y(e_2) \wedge \ell_p(e_1) = \ell_g(e_2) \right\}$$

The yield $y(e) \subseteq W$ of an edge $e = (u, v)$ in either graph is the set of terminals in W that are descendants of v . ℓ is the edge label.

Labeled precision, recall and F-score are then defined as:

$$\text{LP} = \frac{|M(G_p, G_g)|}{|E_p|}, \quad \text{LR} = \frac{|M(G_p, G_g)|}{|E_g|},$$


$$\text{LF} = \frac{2 \cdot \text{LP} \cdot \text{LR}}{\text{LP} + \text{LR}}.$$

Two variants: one for primary edges, and another for remote edges.

Evaluation

Comparing graphs over the same sequence of tokens,

- Match edges by their terminal yield and label.
- Calculate labeled precision, recall and F1 scores.
- Separate primary and remote edges.




Primary: $\frac{LP}{6/9 = 67\%}$ $\frac{LR}{6/10 = 60\%}$ $\frac{LF}{64\%}$

Remote: $\frac{LP}{1/2 = 50\%}$ $\frac{LR}{1/1 = 100\%}$ $\frac{LF}{67\%}$

Bilexical Graph Approximation

No existing UCCA parsers \Rightarrow compare to bilexical parsers:

1. Convert UCCA to bilexical dependencies.
2. Train bilexical parsers and apply to test sentences.
3. Reconstruct UCCA graphs and compare with gold standard.



Bilexical DAG approximation (for tree, delete remote edges).